

# COMMUNICATING RESULTS

*Trevor Lindsay*

---

## COMMUNICATING RESULTS

---

# LEARNING OBJECTIVES

- Explain the trade-offs between the precision and recall of a model while articulating the cost of false positives vs. false negatives.
- Describe the difference between visualization for presentations vs. exploratory data analysis
- Identify the components of a concise, convincing report and how they relate to specific audiences/stakeholders

---

**COURSE**

---

**PRE-WORK**

---

## **PRE-WORK REVIEW**

---

- Understand results from a confusion matrix and measure true positive rate and false positive rate
- Create and interpret results from a binary classification problem

---

**OPENING**

---

# COMMUNICATING RESULTS

---

## **WE BUILT A MODEL...NOW WHAT?**

---

- We've built our model, but there is still a gap between your Notebook with plots/figures and a slideshow needed to present your results.
- Classes so far have focused on two core concepts:
  - developing consistent practices
  - interpreting metrics to evaluate and improve model performance
- But what does that mean to your audience?

---

## WE BUILT A MODEL...NOW WHAT?

---

- Imagine how a non-technical audience might respond to the following statements:
  - The predictive model I built has an accuracy of 80%.
  - Linear regression was optimized by ensuring no multicollinearity
  - Gender was more important than age in the predictive model because it has a larger coefficient.
  - Here's the confusion matrix that shows how well the model did.

---

## WE BUILT A MODEL...NOW WHAT?

---

- Who is your audience? Are they technical? What are their concerns?
- Remember: in a business setting, you may be the only person who can interpret what you've built.
- Some people may be familiar with basic visualization, but you will likely have to do a lot of “hand holding”.
- You need to be able to efficiently explain your results in a way that makes sense to **all** stakeholders (technical or not).



---

## **WE BUILT A MODEL...NOW WHAT?**

---

- Today, we'll focus on communicating results for “simpler” problems, but this applies to any type of model you may work with.
- First, let's review classification metrics, review our knowledge, and talk about how we might communicate what we know.

---

**REVIEW**

---

# BACK TO THE CONFUSION MATRIX

---

## **BACK TO THE CONFUSION MATRIX**

---

- Confusion matrices allow for the interpretation of correct and incorrect predictions for each class label.
- It is the first step for the majority of classification metrics and goes deeper than just accuracy.

# BACK TO THE CONFUSION MATRIX

- Let's recall our confusion matrix

		<u>True class</u>	
		<b>p</b>	<b>n</b>
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives
	<b>N</b>	False Negatives	True Negatives

Column totals:

**P**

**N**

$$\text{fp rate} = \frac{FP}{N}$$

$$\text{tp rate} = \frac{TP}{P}$$

$$\text{precision} = \frac{TP}{TP+FP} \quad \text{recall} = \frac{TP}{P}$$

$$\text{accuracy} = \frac{TP+TN}{P+N}$$

$$\text{F-measure} = \frac{2}{1/\text{precision} + 1/\text{recall}}$$

# THE MATH FOR RECALL

- Recall is the count of predicted true positives over the total count of that class label.
- This is the same as True Positive Rate or sensitivity.


		<u>True class</u>			
		<b>p</b>	<b>n</b>		
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives	$fp\ rate = \frac{FP}{N}$	$tp\ rate = \frac{TP}{P}$
	<b>N</b>	False Negatives	True Negatives	$precision = \frac{TP}{TP+FP}$	$recall = \frac{TP}{P}$
Column totals:		<b>P</b>	<b>N</b>	$accuracy = \frac{TP+TN}{P+N}$	
				$F\text{-measure} = \frac{2}{1/precision + 1/recall}$	

# THE MATH FOR PRECISION

- Precision, or positive predicted value, is calculated as the count of predicted true positives over the count of all values predicted to be positive

		<u>True class</u>			
		<b>p</b>	<b>n</b>		
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives	$fp\ rate = \frac{FP}{N}$	$tp\ rate = \frac{TP}{P}$
	<b>N</b>	False Negatives	True Negatives	$precision = \frac{TP}{TP+FP}$	$recall = \frac{TP}{P}$
Column totals:		<b>P</b>	<b>N</b>	$accuracy = \frac{TP+TN}{P+N}$	
				$F\text{-measure} = \frac{2}{1/precision+1/recall}$	

# VISUALIZING PRECISION AND RECALL



**DEMO**

---

# UNDERSTANDING TRADEOFFS



---

## UNDERSTANDING TRADEOFFS

---

- Let's consider the following data problem: we are given a data set in order to predict or identify traits for typically late flights.
- Optimizing toward recall, we could assume that every flight will be delayed.
- The trade-off, a lower precision, is that this could create even further delays, missed flights, etc.

---


## UNDERSTANDING TRADEOFFS

---

- Optimizing toward precision, we would specifically look to identify flights that will be late.
- The trade-off here would be lower recall. We might miss flights that would be delayed, causing a strain on the system.

# PRECISION RECALL CURVE

- ▶ Another way to visualize a model and compare between models



---

## INTRODUCTION

---

# SHOWING WORK

---

## SHOWING WORK

---

- We've spent a lot of time exploring our data and building a reasonable model that performs well.
- However, if we look at our visuals, they are most likely:
  - Statistically heavy: Most people don't understand histograms.
  - Overly complicated: Scatter matrices produce too much information.
  - Poorly labeled: Code doesn't require adding labels, so you may not have added them.

---

## **SHOWING WORK**

---

- In order to convey important information to our audience, make sure our charts are:
  - Simplified
  - Easily interpretable
  - Clearly labeled

---

## **SIMPLIFIED**

---

- At most, you'll want to include figures that either explain a variable on its own or explain that variable's relationship with a target.
- If your model used a data transformation (like natural log), just visualize the original data.
- Try to remove any unnecessary complexity.

---

## **EASILY INTERPRETABLE**

---

- Any stakeholder looking at a figure should be seeing the exact same thing you're seeing.
- A good test for this is to share the visual with others less familiar with the data and see if they come to the same conclusion.
- How long did it take them?



---

## **CLEARLY LABELED**

---

- Take the time to clearly label your axis, title your plot, and double check your scales - especially if the figures should be comparable.
- If you're showing two graphs side by side, they should follow the same Y axis.

---

## QUESTIONS TO ASK

---

- When building visuals for another audience, ask yourself these questions:
  - **Who:** Who is my target audience for the visual?
  - **What:** What do they already know about this project? What do they need to know?
  - **How:** How does my project affect this audience? How might they interpret (or misinterpret) the data?

**DEMO**

---

# VISUALIZING MODELS OVER VARIABLES

---

## **VISUALIZING MODEL OVER VARIABLES**

---

- One effective way to explain your model over particular variables is to plot the predicted values against the most explanatory variables.
- For example, in logistic regression, plotting the probability of a class against a variable can help explain the range of effect of the model.

---

## **VISUALIZING MODEL OVER VARIABLES**

---

- We'll use the flight delay data for all following examples. Let's build our first model and plot.
- <https://www.rita.dot.gov/bts/help/aviation/index.html>
- Open the a Jupyter notebook and follow along

---

# VISUALIZING MODEL OVER VARIABLES

---

*# read in the file and generate a quick model (assume we've done the data exploration already)*

```
import pandas as pd
```

```
import numpy as np
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
plt.style.use('seaborn')
```

```
import sklearn.linear_model as lm
```

```
df = pd.read_csv('../data/flight_delays.csv')
```

---

# VISUALIZING MODEL OVER VARIABLES

---

*# Remove rows where target, DEP\_DEL15, is null*

```
df = df[df['DEP_DEL15'].notnull()].copy()
```

*# Convert UNIQUE\_CARRIER to a dummy variable*

```
df = pd.get_dummies(df, prefix='CARRIER', columns=['UNIQUE_CARRIER'])
```

*# Convert DAY\_OF\_WEEK to a dummy variable*

```
df = pd.get_dummies(df, prefix='DOW', columns=['DAY_OF_WEEK'], drop_first=True)
```

---

# VISUALIZING MODEL OVER VARIABLES

---

*# Create features and target set*

```
features = [col for col in df.columns if 'DOW_' in col] + ['CRS_DEP_TIME']
```

```
target = 'DEP_DEL15'
```

```
X = df[features]
```

```
y = df[target]
```

*# Build a model*

```
cls = lm.LogisticRegression()
```

```
cls.fit(X, y)
```

*# Predict probability of delay*

```
df['probability'] = cls.predict_proba(X)[:, 1]
```



---


# VISUALIZING MODEL OVER VARIABLES

---

```
ax = plt.gca()
colors = ['blue', 'green', 'red', 'purple', 'orange', 'brown']
for index, color in enumerate(colors):
    dow = features[index]
    df[df[dow] == 1].plot(
        x='CRS_DEP_TIME',
        y='probability',
        kind='scatter',
        color=color,
        figsize=(10,8),
        ax=ax,
    )
ax.set(title='Probability of Delay\n Based on Day of Week and Time of Day')
```

# VISUALIZING MODEL OVER VARIABLES

- This visual can help showcase the range of effect on delays from both day of the week and time of day.
- Given this model, some days are more likely to have delays than others.
- The likelihood of delay increases as the day goes on.



---

## ACTIVITY: TRY IT OUT

---

- Directions

- 1. Adjust the model to make delay predictions using airlines instead of day of week, and time, then plot the effect on `CRS_DEP_TIME=1`
  - 2. Try plotting the inverse: pick either model and plot the effect on `CRS_DEP_TIME=0`

- Deliverable

- The new plots

**DEMO**

---

# **VISUALIZING PERFORMANCE AGAINST BASELINE**

---


## **VISUALIZING PERFORMANCE AGAINST BASELINE**

---

- Another approach of visualization is the effect of your model against a baseline, or -- even better -- against previous models.
- Plots like this will also be useful when talking to your peers - other data scientists or analysts who are familiar with your project and interested in the progress you've made.


# VISUALIZING PERFORMANCE AGAINST BASELINE

- For classification, we've practiced plotting AUC and precision-recall plots. Consider the premise of each:
- AUC plots explain and represent “accuracy” as having the largest area under the curve. Good models will be high and to the left.
- For precision-recall plots, it will depend on the cost requirements. Either a model will have good recall at the cost of precision or vice versa.



# VISUALIZING PERFORMANCE AGAINST BASELINE

- ▶ When comparing multiple models:
- ▶ For AUC plots, you'll be interested in which model has the largest area under the curve.
- ▶ For precision-recall plots, based on the cost requirement, you are looking at which model has the best precision given the same recall, or the best recall given the same precision.



---

# VISUALIZING PERFORMANCE OVER BASELINE

---

```
from sklearn.dummy import DummyClassifier
```

```
from sklearn import metrics
```

```
dummy = DummyClassifier()
```

```
dummy.fit(X, y)
```

```
df['probability_dummy'] = dummy.predict_proba(X)[:, 1]
```



---

# VISUALIZING PERFORMANCE OVER BASELINE

---

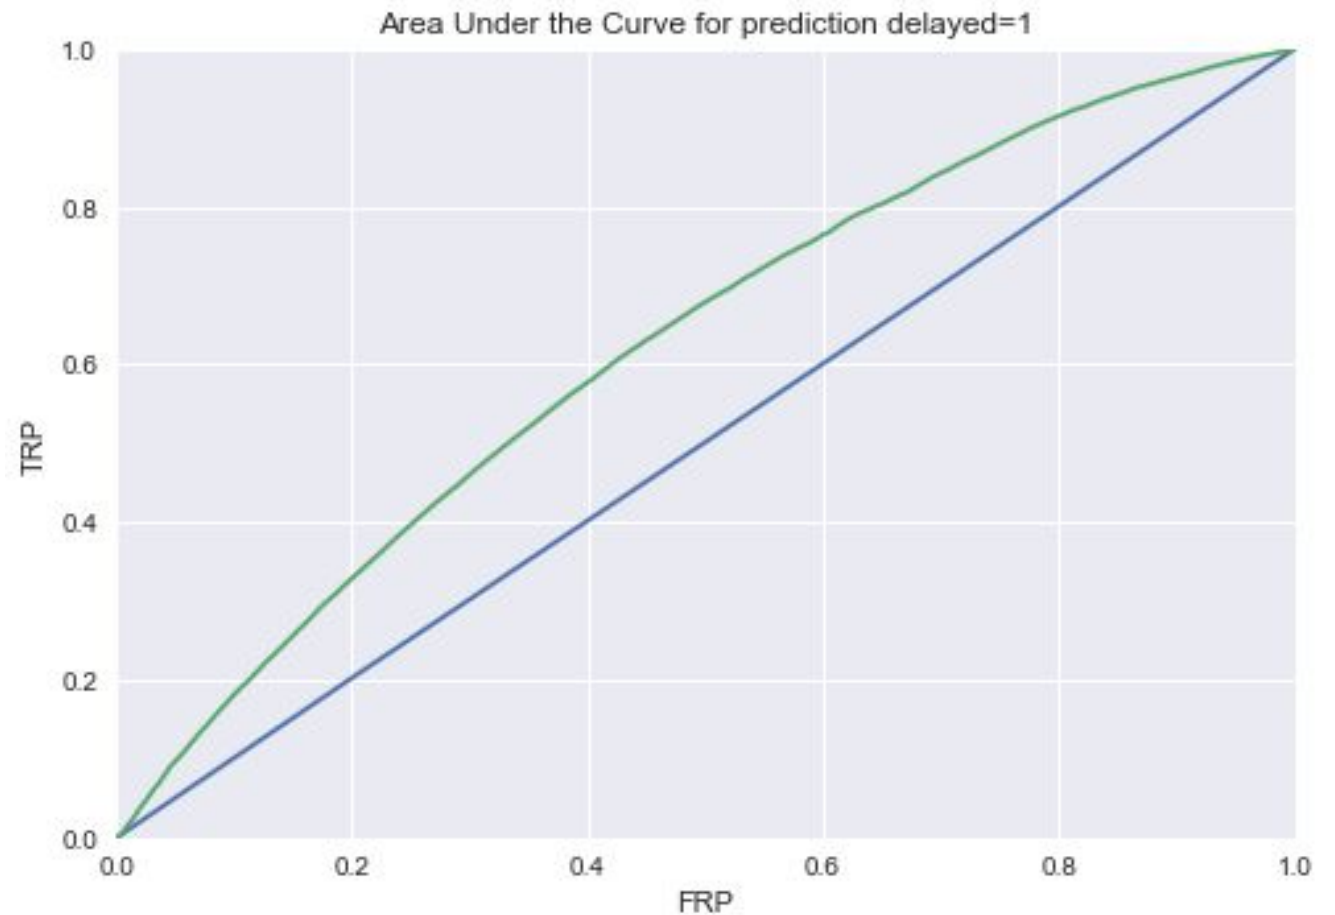
```
ax = plt.gca()

vals = metrics.roc_curve(y, df['probability_dummy'])
ax.plot(vals[0], vals[1])
vals = metrics.roc_curve(y, df['probability'])
ax.plot(vals[0], vals[1])

ax.set(
    title='Area Under the Curve for prediction delayed=1',
    ylabel='TRP',
    xlabel='FRP',
    xlim=(0, 1),
    ylim=(0, 1)
)
```

# VISUALIZING PERFORMANCE AGAINST BASELINE

- This plot showcases:
  - The model using data outperforms a baseline dummy model.
  - By adding other features, there's some give and take with probability as the model gets more complicated.



---

## ACTIVITY -- TRY IT OUT

---

- In a similar approach, use the `sklearn precision_recall_curve` function to enable you to plot the precision-recall curve of the four models from above.
  - Keep in mind precision in the first array is returned from the function, but the plot shows it as the y-axis
- 2. Explain what is occurring when the recall is below 0.2.
- 3. Based on this performance, is there a clear winner at different thresholds?
- 4. **Bonus:** Redo both the AUC and precision-recall curves using models that have been cross validated using `kfold`. How do these new figures change your expectations for performance?

---

**INDEPENDENT PRACTICE**

---

# PROJECT PRACTICE

---

## ACTIVITY -- TRY IT OUT

---

- Using models built from the flight data problem earlier in class, work through the same problems. Your data and models should already be accessible.
- Your goals:
  - There are many ways to manipulate this data set. Perhaps mix both carrier and day of week in your model. Consider what is a proper "categorical" variable, and keep only what is significant. Aim to have a visual that clearly explain the relationship of variables you've used against the predicted flight delay.
  - Generate the AUC or precision-recall curve (based on which you think makes more sense), and have a statement that defines, compared to a baseline, how your model performs and any caveats. For example: "My model on average performs at x rate, but the features under-perform and explain less of the data at these thresholds."

---

**REVIEW**

---

# TOPIC REVIEW

---

## TOPIC REVIEW

---

- What do precision and recall mean? How are they similar and different to True Positive Rate and False Positive Rate?
- How does cost benefit analysis play a role in building models?
- What are at least two very important details to consider when creating visuals for a project's stakeholders?
- Why would an AUC plot work well for a data science audience but not for a business audience? What would be a more effective visualization for that group?